

# Ajouter une autre dimension a la qualité des logiciels

Exemple de mise en oeuvre d'un outil  
d'analyse statique

Rene Brun, CERN

# CERN

- Organisation Européenne pour la Recherche Nucléaire (Physique des particules): Frontière Franco-Suisse près de Genève
- Recherche fondamentale, mais aussi avec des retombées pratiques (électronique, vide, cryogénie, scanners et le WEB)
- Comprendre le big-bang et l'évolution de l'univers

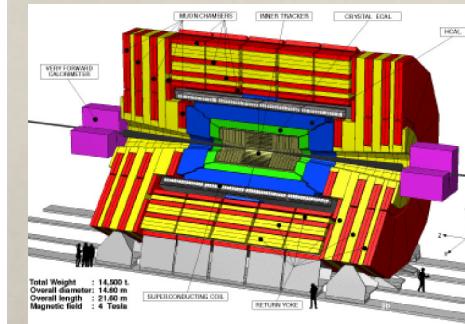
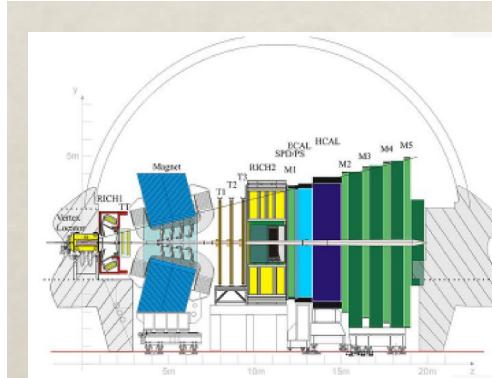
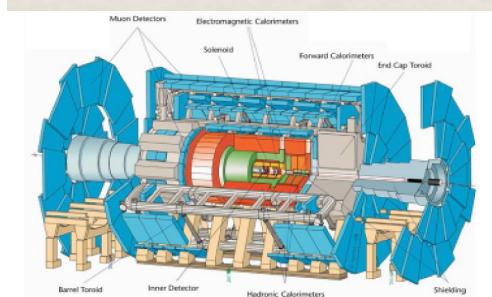


# Le LHC (Large Hadron Collider)

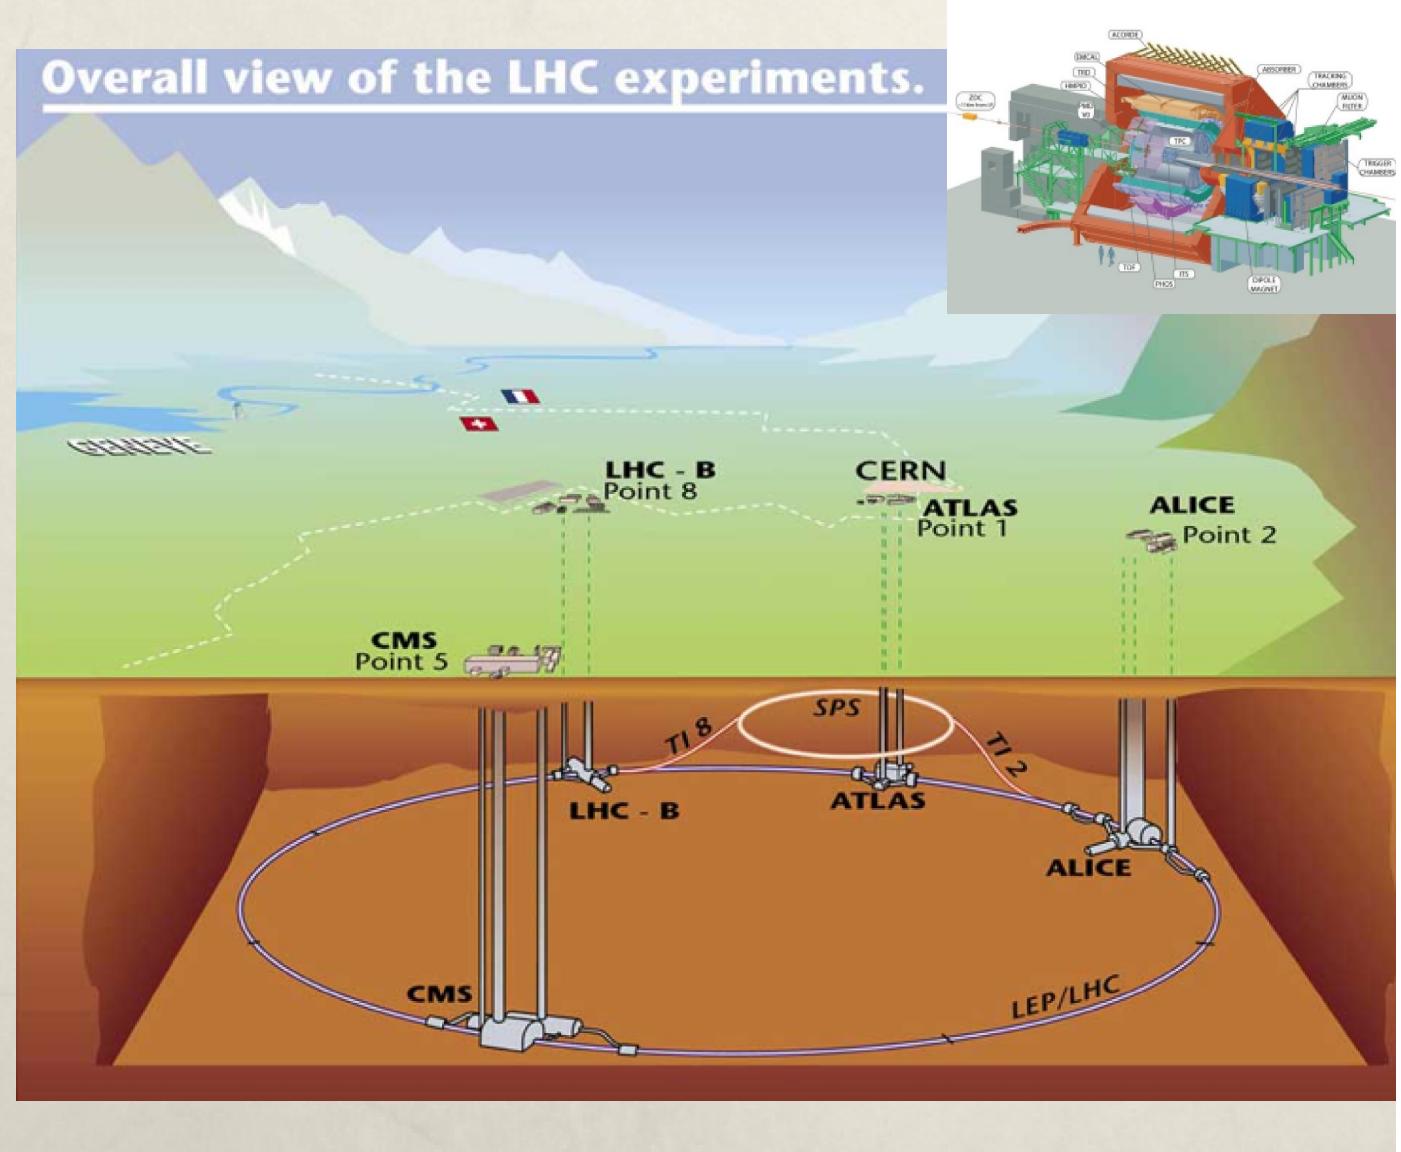
- Un tunnel de 28km a 100m sous terre a la frontiere FR/CH
- Collisioneur de protons ou d'ions lourds
- 10000 physiciens



# LHC



## Overall view of the LHC experiments.



# CERN en chiffres

- Utilisé par 10000 physiciens de 608 universités dans 113 pays.
- Plus de 1000 physiciens par expérience
- Une expérience peut durer 40 ans!
- The LHC génère 15,000,000 Gigabytes de données par an.
- Budget: 1 milliard de CHF/an

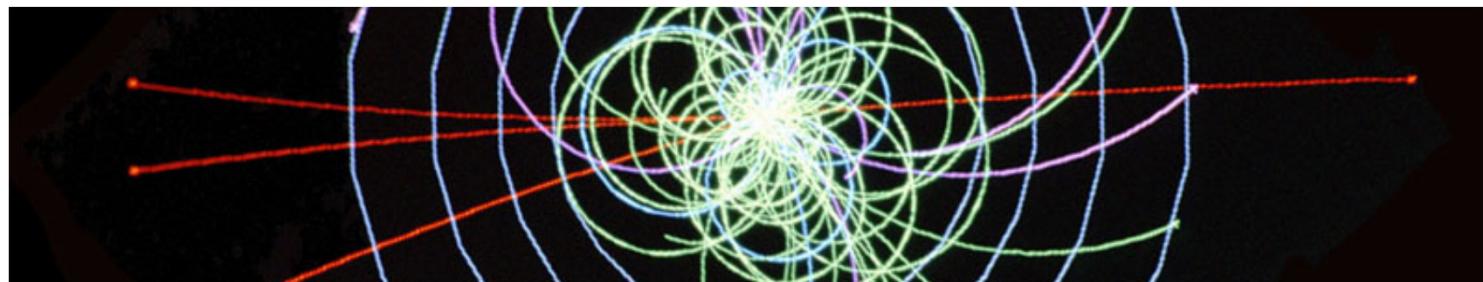
# Informatique au CERN

- Approx 50 MLOC C++ par les physiciens
- Beaucoup de données: code doit être rapide
- The code doit être correct: résultats *scientifiques*;
- Doit être stable: 1 Higgs / 10,000,000,000,000 collisions;



# Logiciels au CERN

- Chaque expérience développe et maintient ses propres logiciels (2/3 du total environ).
- The CERN développe et maintient des outils généralistes communs à toutes les expériences: ROOT: math / statistiques, stockage et management des données, simulation de détecteurs, graphiques, GUI, calculs en parallèle.



# Logiciels des expériences

- Plusieurs centaines de développeurs/expérience: physiciens, quelques experts logiciels; actifs pour quelques années seulement (beaucoup de PostDoc).
- Développement très souvent anarchique.
- Beaucoup de dépendances sur des logiciels externes.

# ROOT

- <http://root.cern.ch>
- 7 Developpeurs @ 2.6 MLOC
- Début en 1995, 90% du code avant 2005. Doit être maintenu sur un grand nombre de systèmes durant au moins 30 ans.
- Deux releases / an
- “Core project”, utilise par toutes les expériences/ physiciens y compris en dehors du domaine du CERN: 95000 adresses IP distinctes/an.

# Software Risks

- Maintenance
- Implosion of huge software stack
- Integrity of years of LHC experiments' data
- Correctness of scientific results
- Failing searches for new physics
- *Note:* driven by reality, not ISO xyz

# Challenges For QA

- C++
- anarchy
- few experts
- multi-platform
- huge programs
- part-time coding
- input data volume
- software as a tool
- interpreter binding
- lack of reproducibility

# QA Requirements

- Results must be easy to understand: novices, spare time coding, complex code
- Results must be relevant: motivate, don't flood, developers
- QA must handle rare code paths: don't hide the Higgs
- QA as part of nightly / continuous testing

# QA Toolbox

- Patch review (ROOT)
- Continuous integration (ROOT)
- Multi-Platform: compilers, versions, architectures
- Unit + functionality test suites 100k LOC for ROOT
- Regression tests on performance, results

# QA Toolbox (2)

- Nightly builds + tests + automatic reports
- Memory: valgrind, massif, ROOT, igprof...
- Performance: callgrind, shark, gprof...
- Test coverage: gcov
- Coding rule checkers
- Automation + strong software management

# Memory QA

- C++ inherent: memory errors
  - use after delete
  - uninitialized pointers
  - buffer overruns
  - memory leaks
- Amplified by novice coders, non-defined ownership rules,...

# Performance QA

- Huge amounts of data, limited computing, thus performance sensitive
- Difficult to predict: caching, templates, inlining, call of external libraries
- Monitoring difficult to present, interpret
- Using advanced tools + techniques to handle >10 MLOC in a novice-friendly way

# QA Toolbox: Full?

```
void f(int* p) {  
    g(p);  
    if (p) *p = 12;  
}
```

```
void g(int* p) {  
    *p = 0;  
}
```

```
char buf[1024];  
strcpy(buf, getenv("PATH"));
```

```
void p(int flag) {  
    if (flag > 2) {  
        ...  
        return;  
    }  
    int flag1 = flag * 2;  
    if (flag1 < 10) {  
        ...  
    } else {  
        // Algorithm that a  
        // physicist worked  
        // on for two years  
    }
```

```
if ( a == '1' || '2' ) {
```

```
void receive() {  
    if (!m_Server->receive())  
        recover(S);  
    logger(S->Name(), "receive");  
}
```

```
void recover(Server* S) {  
    cout << "ERROR!\n";  
    delete S;  
}
```

# Open QA Issues

- Those *almost* impossible code paths!
- Problems that cannot be reproduced
- Too many code paths for humans
- Condition matrices
- “You know what I mean” coding

# Static Code Checkers

- Check what *could be* run, not what *is* run
  - CERN did a study of open source options
- 
- /media/thomas/userprovideddata/mato/root/build/rmkdepend/include.c:295: [4] (buffer)  
*strcpy: Does not check for buffer overflows when copying to destination. Consider using strncpy or strlcpy (warning, strncpy is easily misused).*
  - /media/thomas/userprovideddata/mato/root/build/rmkdepend/include.c:316: [4] (buffer)  
*sprintf: Does not check for buffer overflows. Use snprintf or vsnprintf.*
  - /media/thomas/userprovideddata/mato/root/build/rmkdepend/main.c:553: [4] (buffer)  
*strcpy: Does not check for buffer overflows when copying to destination. Consider using strncpy or strlcpy (warning, strncpy is easily misused).*
  - /media/thomas/userprovideddata/mato/root/build/rmkdepend/main.c:704: [4] (buffer)  
*sprintf: Does not check for buffer overflows. Use snprintf or vsnprintf.*
  - /media/thomas/userprovideddata/mato/root/build/rmkdepend/main.c:752: [4] (format)  
*vfprintf: If format strings can be influenced by an attacker, they can be exploited. Use a*

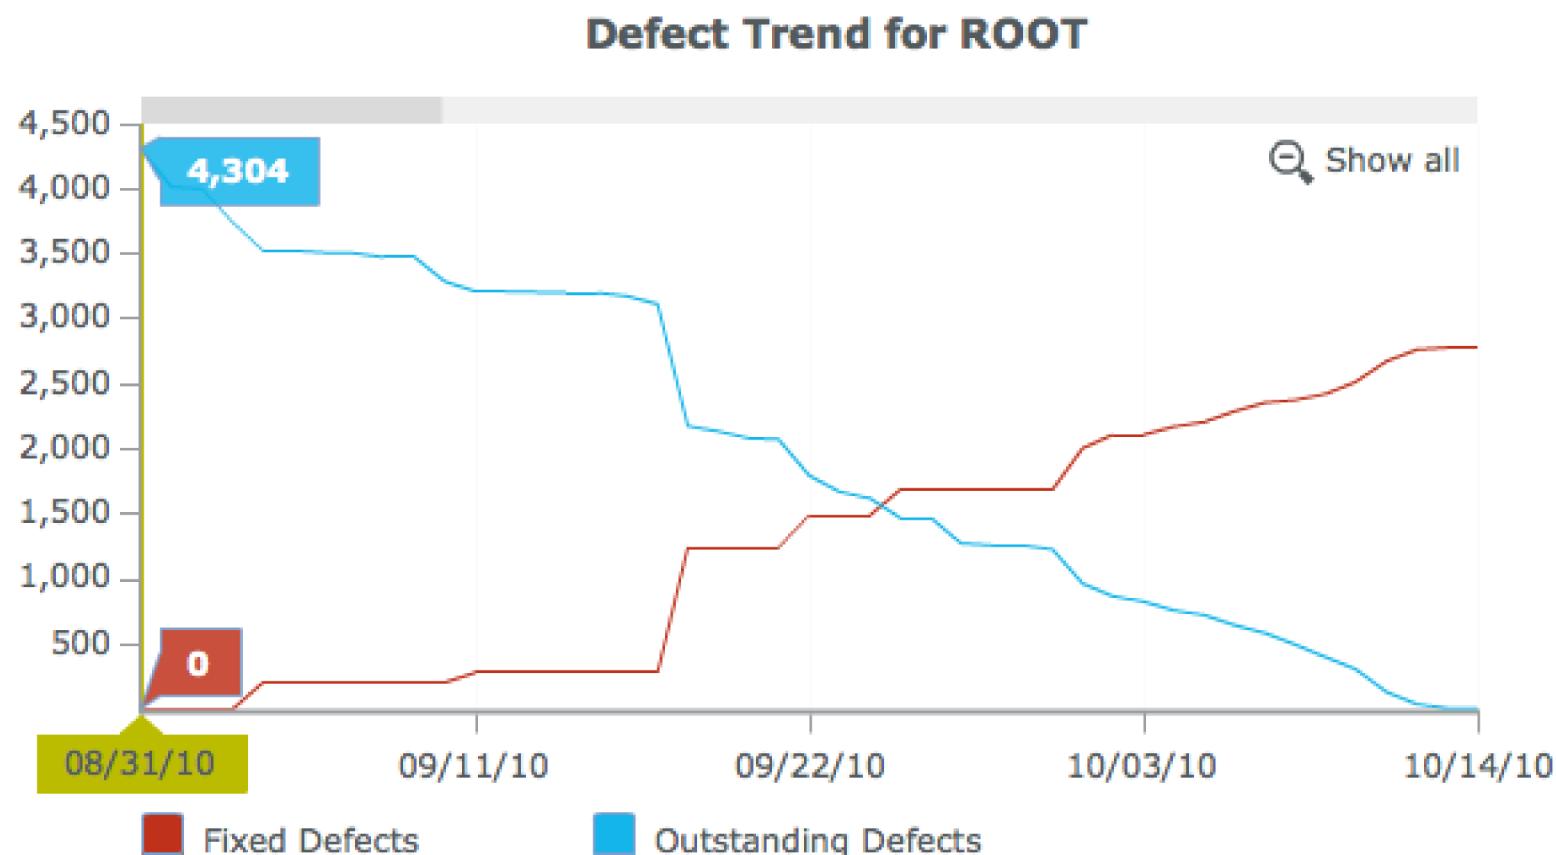
# Coverity @ CERN

- ROOT traditionally (informally) spearheads new technologies at CERN: what works for ROOT (and scales!) works for CERN
- Realized lack of static analysis at CERN
- Discovered Coverity through famous Scan: Linux, Python, Apache, glibc, openSSL,...
- ROOT as CERN's guinea pig project

# ROOT Meets Coverity

- Plan:
  - Product discovered by a team member
  - Convince ROOT developers
  - Convince CERN
- Set up as “idle time project” in 4 weeks: build integration, automation, reporting
- 2.6 MLOC: 4 processes \* 8 hours

# First Coverity Results

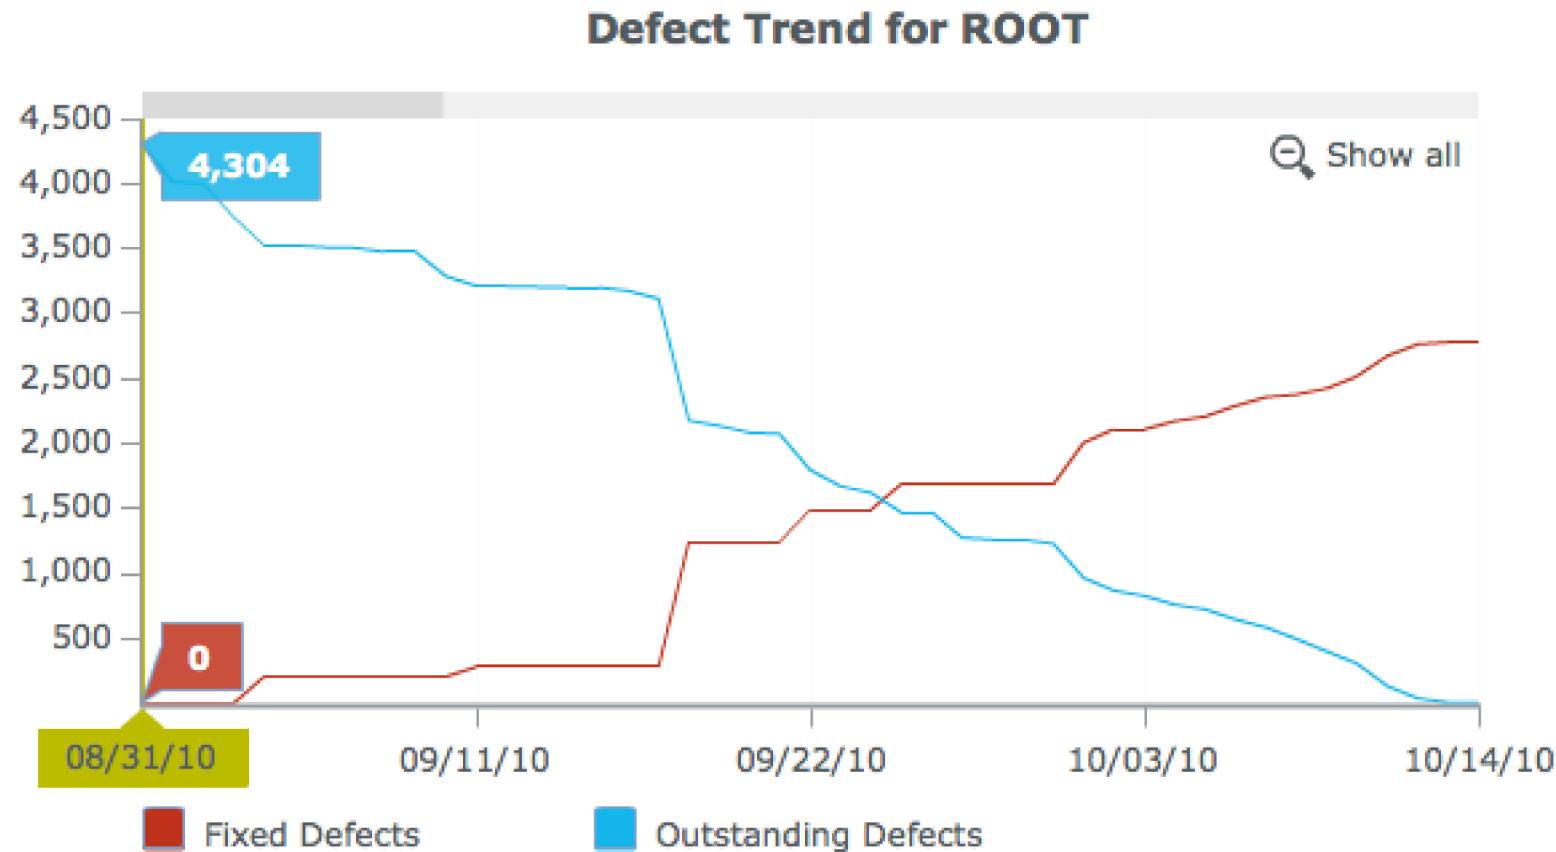


# Report Quality

- Very readable reports embedded in source
- Mostly easy to understand, though complex contexts make complex reports
- Amazing Relevance: 12% false positives, 23% intentional thus 2/3 reports causing source changes - despite an already well-filled QA toolbox

# Consequence

- Relevant reports create motivated developers!



# Results With ROOT

- 4300 reports by 7 developers in 6 weeks
- Noticeably better release
  - Instability due blind changes, e.g. “missing break”
- Found several long-hunted bugs, e.g. hard to reproduce or track down
- Systematic deficiencies triggered systematic fixes: buffer overflows / strlcpy

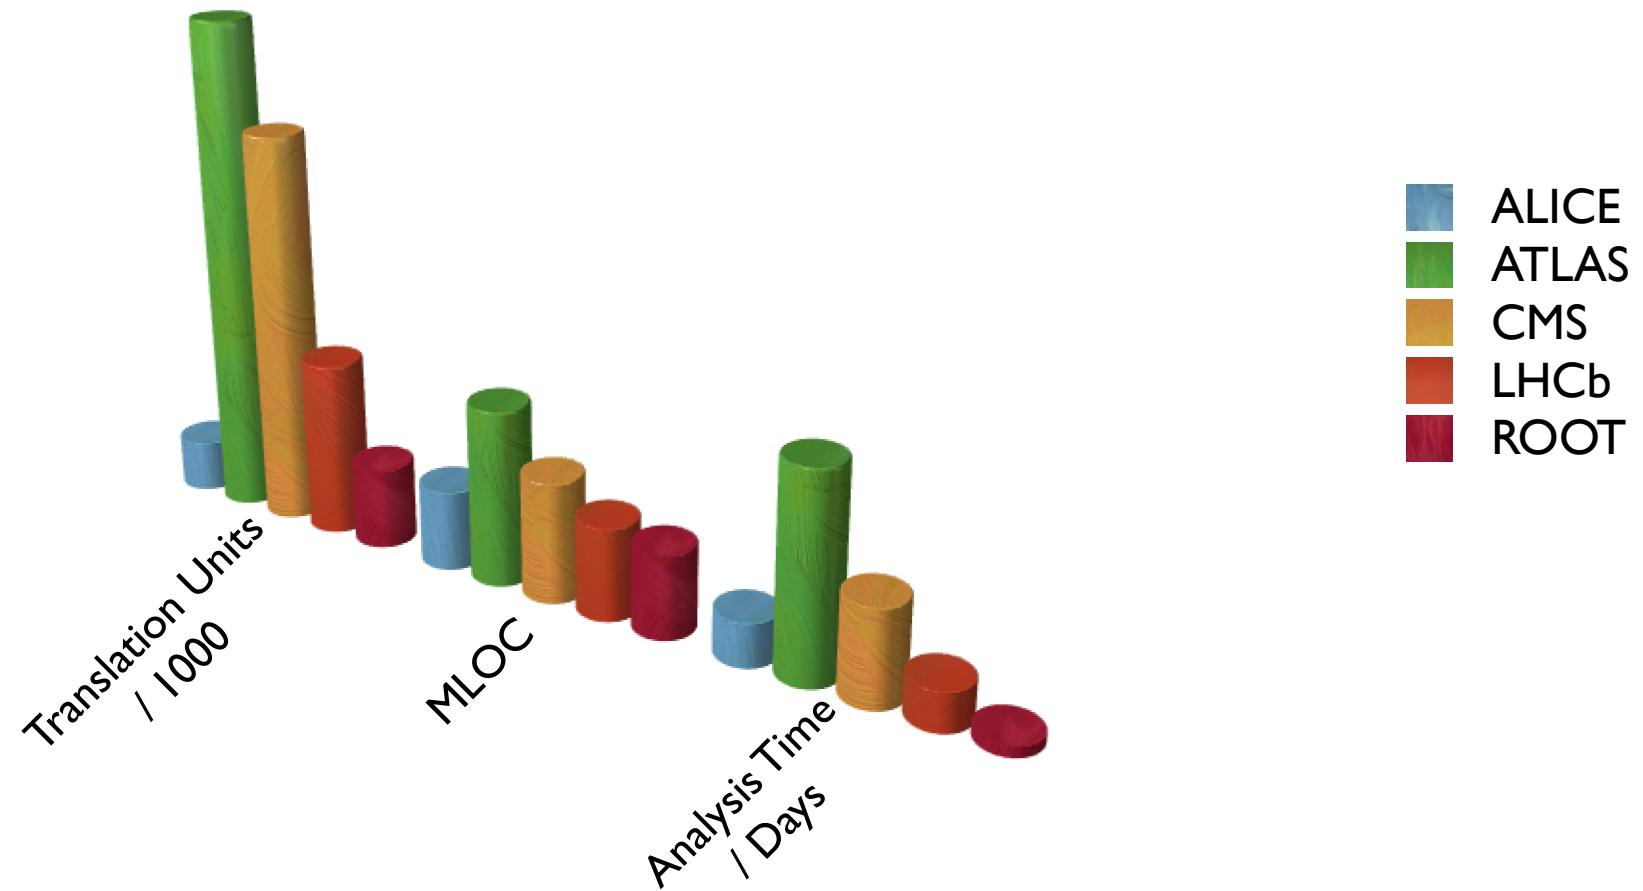
# Coverity On 50MLOC

- Excellent show-case for experiments
- Immediate interest:
  - Allows for precise, relevant, accessible, large-scale, automated, quick turn-around reports for software management
  - Classification in high / low impact helps fixing critical reports first
  - Increased awareness of software quality issues

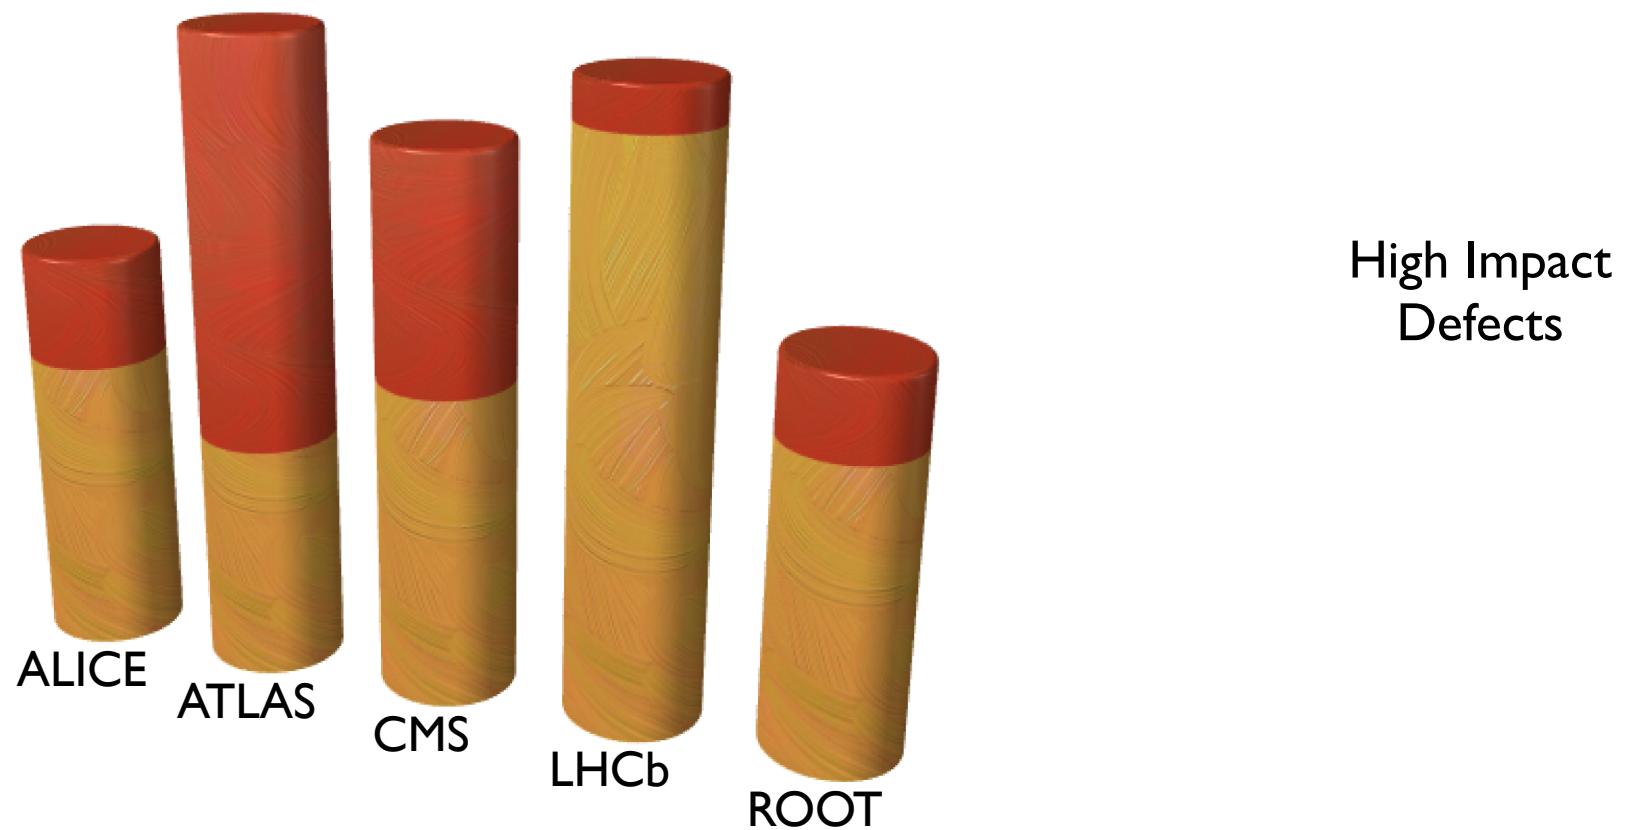
# Setup

- Separate build + analysis server for each experiment due to processing times: 4 processes \* 8 hours to 4 days - partially caused by build systems
- Dedicated database / tomcat server for web front-end, for each experiment
- Keeps databases reasonably small, user (web) interface responsive

# Initial Numbers: Time



# Initial Numbers: Issues



# Has It Arrived?

- ROOT, particle beam, data transfer, virtual machine etc: nightly / regular use
- All LHC experiments use it, one even embedded in automatic QA
- Coverity is known in all experiments as “the tool that doesn’t forgive lazy coding”

# Summary

- Software quality a key factor for CERN: influences quality of scientific results
- Blind spot despite of large QA toolset
- Coverity unique in abilities and quality
  - Improved software, happier users
  - Quick adoption at CERN (rare!): recognized to improve software quality dramatically